

SyriaTel Customer Churn Prediction Project

Business Understanding

Business Problem: Reducing the rate of customer churn for SyriaTel Telecom

Background: SyriaTel, a major player in the telecommunications industry, is dealing with a serious problem of customer churn, which occurs when customers cancel their subscriptions, leading to a loss of revenue. To tackle this challenge, SyriaTel plans to develop a predictive model to identify customers who are at risk of churning. By reaching out to these customers with targeted retention strategies, SyriaTel aims to decrease churn rates and keep valuable customers.

Problem statement: Can we predict customer churn for SyriaTel and determine the main factors contributing to it, allowing the company to develop effective retention strategies.

Objectives:

- Develop a Churn Prediction model capable of accurately identifying customers likely to churn using the available dataset.
- Discover the key patterns and features that contribute to predicting customer churn.

This initiative supports SyriaTel's objective of retaining customers and minimizing revenue loss from churn, highlighting the value of data-driven decision-making in the telecommunications sector.

Expected Outcome

The main metric for evaluating the classification model's performance is 'recall,' which measures how well the model identifies customers likely to churn. The goal is to minimize false negatives since missing a potential churner is more costly for the business than incorrectly predicting a non-churner. The target is to achieve at least 80% recall.

However, a balance is necessary. Predicting that all customers will churn would result in perfect recall but offer no real business value, as not all customers are at risk. Therefore, in addition to recall, 'precision' and 'accuracy' will be used as secondary metrics to ensure a comprehensive assessment of the model's performance.

Import the libraries

Here, we import the various necessary libraries that we will be using in this project. They include Pandas, Numpy, visualization libraries like matplotlib and seaborn, various scikit-learn modules for machine learning and metrics we will use for evaluating our models.

```
In [35]: # Importing the necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import math
%matplotlib inline

import scipy.stats as stats

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from imblearn.over_sampling import SMOTE
from sklearn.svm import SVC
from sklearn.dummy import DummyClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrix
from sklearn.feature_selection import RFECV

import warnings
warnings.filterwarnings('ignore')
```

Data Understanding

In the data folder is a "telcom_data.csv" file, a dataset from SyriaTel Telecom, that contains data about various customers, and whether they churned or not.

```
In [36]: pd.set_option('display.max_columns', None)
```

```
In [37]: # Loading the dataset
df = pd.read_csv('C:/Users/HP/AppData/Local/Temp/81e1e7e9-e08e-41c6-a8e3-d22ea09a02d5')
df.head()
```

```
Out[37]:
```

| | state | account length | area code | phone number | international plan | voice mail plan | number vmail messages | total day minutes | total day calls | total day charge | total eve minutes |
|---|-------|-------------------|--------------|-----------------|-----------------------|-----------------------|-----------------------------|-------------------------|-----------------------|------------------------|-------------------------|
| 0 | KS | 128 | 415 | 382-4657 | no | yes | 25 | 265.1 | 110 | 45.07 | 197.4 |
| 1 | OH | 107 | 415 | 371-7191 | no | yes | 26 | 161.6 | 123 | 27.47 | 195.5 |
| 2 | NJ | 137 | 415 | 358-1921 | no | no | 0 | 243.4 | 114 | 41.38 | 121.2 |
| 3 | OH | 84 | 408 | 375-9999 | yes | no | 0 | 299.4 | 71 | 50.90 | 61.9 |

| | state | account length | area code | phone number | international plan | voice mail plan | number vmail messages | total day minutes | total day calls | total day charge | total eve minutes |
|---|-------|----------------|-----------|--------------|--------------------|-----------------|-----------------------|-------------------|-----------------|------------------|-------------------|
| 4 | OK | 75 | 415 | 330-6626 | yes | no | 0 | 166.7 | 113 | 28.34 | 148.3 |

```
In [38]: # Checking last 4 records
df.tail()
```

```
Out[38]:
```

| | state | account length | area code | phone number | international plan | voice mail plan | number vmail messages | total day minutes | total day calls | total day charge | total eve minutes |
|------|-------|----------------|-----------|--------------|--------------------|-----------------|-----------------------|-------------------|-----------------|------------------|-------------------|
| 3328 | AZ | 192 | 415 | 414-4276 | no | yes | 36 | 156.2 | 77 | 26.55 | 215 |
| 3329 | WV | 68 | 415 | 370-3271 | no | no | 0 | 231.1 | 57 | 39.29 | 153 |
| 3330 | RI | 28 | 510 | 328-8230 | no | no | 0 | 180.8 | 109 | 30.74 | 286 |
| 3331 | CT | 184 | 510 | 364-6381 | yes | no | 0 | 213.8 | 105 | 36.35 | 159 |
| 3332 | TN | 74 | 415 | 400-4344 | no | yes | 25 | 234.4 | 113 | 39.85 | 265 |

From the above preview of the dataset, we can see that the dataset is uniform.

```
In [39]: # Checking the shape of the dataset
df.shape
```

```
Out[39]: (3333, 21)
```

```
In [40]: # Overview of the columns
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                        3333 non-null   object
4   international plan                   3333 non-null   object
5   voice mail plan                      3333 non-null   object
6   number vmail messages                3333 non-null   int64
7   total day minutes                    3333 non-null   float64
8   total day calls                      3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                    3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
18  total intl charge                    3333 non-null   float64
19  customer service calls               3333 non-null   int64
20  churn                               3333 non-null   bool
```

```
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

From the above, we can see that the columns are of different types: float, int, object and bool. However, most of the columns have some whitespace in between their names, and it would be better to rename them in such a way that they have an underscore instead of the whitespace.

Renaming the by replacing the whitespace with an underscore and displaying unique values

```
In [41]: df.columns = df.columns.str.replace(' ', '_')

# Display unique values for each column
for column in df.columns:
    print(f"There are {df[column].nunique()} Unique values in '{column}':")
    print(df[column].unique())
    print('-' * 50)
```

There are 51 Unique values in 'state':

```
['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO' 'LA' 'WV' 'IN' 'RI' 'IA' 'MT' 'NY'
 'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
 'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
 'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND']
```

There are 212 Unique values in 'account_length':

```
[128 107 137 84 75 118 121 147 117 141 65 74 168 95 62 161 85 93
 76 73 77 130 111 132 174 57 54 20 49 142 172 12 72 36 78 136
149 98 135 34 160 64 59 119 97 52 60 10 96 87 81 68 125 116
 38 40 43 113 126 150 138 162 90 50 82 144 46 70 55 106 94 155
 80 104 99 120 108 122 157 103 63 112 41 193 61 92 131 163 91 127
110 140 83 145 56 151 139 6 115 146 185 148 32 25 179 67 19 170
164 51 208 53 105 66 86 35 88 123 45 100 215 22 33 114 24 101
143 48 71 167 89 199 166 158 196 209 16 39 173 129 44 79 31 124
 37 159 194 154 21 133 224 58 11 109 102 165 18 30 176 47 190 152
 26 69 186 171 28 153 169 13 27 3 42 189 156 134 243 23 1 205
200 5 9 178 181 182 217 177 210 29 180 2 17 7 212 232 192 195
197 225 184 191 201 15 183 202 8 175 4 188 204 221]
```

There are 3 Unique values in 'area_code':

```
[415 408 510]
```

There are 3333 Unique values in 'phone_number':

```
['382-4657' '371-7191' '358-1921' ... '328-8230' '364-6381' '400-4344']
```

There are 2 Unique values in 'international_plan':

```
['no' 'yes']
```

There are 2 Unique values in 'voice_mail_plan':

```
['yes' 'no']
```

There are 46 Unique values in 'number_vmail_messages':

```
[25 26 0 24 37 27 33 39 30 41 28 34 46 29 35 21 32 42 36 22 23 43 31 38
 40 48 18 17 45 16 20 14 19 51 15 11 12 47 8 44 49 4 10 13 50 9]
```

There are 1667 Unique values in 'total_day_minutes':

```
[265.1 161.6 243.4 ... 321.1 231.1 180.8]
```

There are 119 Unique values in 'total_day_calls':

```
[110 123 114 71 113 98 88 79 97 84 137 127 96 70 67 139 66 90
117 89 112 103 86 76 115 73 109 95 105 121 118 94 80 128 64 106
102 85 82 77 120 133 135 108 57 83 129 91 92 74 93 101 146 72
 99 104 125 61 100 87 131 65 124 119 52 68 107 47 116 151 126 122]
```

```
111 145 78 136 140 148 81 55 69 158 134 130 63 53 75 141 163 59
132 138 54 58 62 144 143 147 36 40 150 56 51 165 30 48 60 42
0 45 160 149 152 142 156 35 49 157 44]
```

```
-----
There are 1667 Unique values in 'total_day_charge':
[45.07 27.47 41.38 ... 54.59 39.29 30.74]
```

```
-----
There are 1611 Unique values in 'total_eve_minutes':
[197.4 195.5 121.2 ... 153.4 288.8 265.9]
```

```
-----
There are 123 Unique values in 'total_eve_calls':
[ 99 103 110 88 122 101 108 94 80 111 83 148 71 75 76 97 90 65
 93 121 102 72 112 100 84 109 63 107 115 119 116 92 85 98 118 74
117 58 96 66 67 62 77 164 126 142 64 104 79 95 86 105 81 113
106 59 48 82 87 123 114 140 128 60 78 125 91 46 138 129 89 133
136 57 135 139 51 70 151 137 134 73 152 168 68 120 69 127 132 143
61 124 42 54 131 52 149 56 37 130 49 146 147 55 12 50 157 155
45 144 36 156 53 141 44 153 154 150 43 0 145 159 170]
```

```
-----
There are 1440 Unique values in 'total_eve_charge':
[16.78 16.62 10.3 ... 13.04 24.55 22.6 ]
```

```
-----
There are 1591 Unique values in 'total_night_minutes':
[244.7 254.4 162.6 ... 280.9 120.1 279.1]
```

```
-----
There are 120 Unique values in 'total_night_calls':
[ 91 103 104 89 121 118 96 90 97 111 94 128 115 99 75 108 74 133
64 78 105 68 102 148 98 116 71 109 107 135 92 86 127 79 87 129
57 77 95 54 106 53 67 139 60 100 61 73 113 76 119 88 84 62
137 72 142 114 126 122 81 123 117 82 80 120 130 134 59 112 132 110
101 150 69 131 83 93 124 136 125 66 143 58 55 85 56 70 46 42
152 44 145 50 153 49 175 63 138 154 140 141 146 65 51 151 158 155
157 147 144 149 166 52 33 156 38 36 48 164]
```

```
-----
There are 933 Unique values in 'total_night_charge':
[11.01 11.45 7.32 8.86 8.41 9.18 9.57 9.53 9.71 14.69 9.4 8.82
6.35 8.65 9.14 7.23 4.02 5.83 7.46 8.68 9.43 8.18 8.53 10.67
11.28 8.22 4.59 8.17 8.04 11.27 11.08 13.2 12.61 9.61 6.88 5.82
10.25 4.58 8.47 8.45 5.5 14.02 8.03 11.94 7.34 6.06 10.9 6.44
3.18 10.66 11.21 12.73 10.28 12.16 6.34 8.15 5.84 8.52 7.5 7.48
6.21 11.95 7.15 9.63 7.1 6.91 6.69 13.29 11.46 7.76 6.86 8.16
12.15 7.79 7.99 10.29 10.08 12.53 7.91 10.02 8.61 14.54 8.21 9.09
4.93 11.39 11.88 5.75 7.83 8.59 7.52 12.38 7.21 5.81 8.1 11.04
11.19 8.55 8.42 9.76 9.87 10.86 5.36 10.03 11.15 9.51 6.22 2.59
7.65 6.45 9. 6.4 9.94 5.08 10.23 11.36 6.97 10.16 7.88 11.91
6.61 11.55 11.76 9.27 9.29 11.12 10.69 8.8 11.85 7.14 8.71 11.42
4.94 9.02 11.22 4.97 9.15 5.45 7.27 12.91 7.75 13.46 6.32 12.13
11.97 6.93 11.66 7.42 6.19 11.41 10.33 10.65 11.92 4.77 4.38 7.41
12.1 7.69 8.78 9.36 9.05 12.7 6.16 6.05 10.85 8.93 3.48 10.4
5.05 10.71 9.37 6.75 8.12 11.77 11.49 11.06 11.25 11.03 10.82 8.91
8.57 8.09 10.05 11.7 10.17 8.74 5.51 11.11 3.29 10.13 6.8 8.49
9.55 11.02 9.91 7.84 10.62 9.97 3.44 7.35 9.79 8.89 8.14 6.94
10.49 10.57 10.2 6.29 8.79 10.04 12.41 15.97 9.1 11.78 12.75 11.07
12.56 8.63 8.02 10.42 8.7 9.98 7.62 8.33 6.59 13.12 10.46 6.63
8.32 9.04 9.28 10.76 9.64 11.44 6.48 10.81 12.66 11.34 8.75 13.05
11.48 14.04 13.47 5.63 6.6 9.72 11.68 6.41 9.32 12.95 13.37 9.62
6.03 8.25 8.26 11.96 9.9 9.23 5.58 7.22 6.64 12.29 12.93 11.32
6.85 8.88 7.03 8.48 3.59 5.86 6.23 7.61 7.66 13.63 7.9 11.82
7.47 6.08 8.4 5.74 10.94 10.35 10.68 4.34 8.73 5.14 8.24 9.99
13.93 8.64 11.43 5.79 9.2 10.14 12.11 7.53 12.46 8.46 8.95 9.84
10.8 11.23 10.15 9.21 14.46 6.67 12.83 9.66 9.59 10.48 8.36 4.84
10.54 8.39 7.43 9.06 8.94 11.13 8.87 8.5 7.6 10.73 9.56 10.77]
```

```
7.73 3.47 11.86 8.11 9.78 9.42 9.65 7. 7.39 9.88 6.56 5.92
6.95 15.71 8.06 4.86 7.8 8.58 10.06 5.21 6.92 6.15 13.49 9.38
12.62 12.26 8.19 11.65 11.62 10.83 7.92 7.33 13.01 13.26 12.22 11.58
5.97 10.99 8.38 9.17 8.08 5.71 3.41 12.63 11.79 12.96 7.64 6.58
10.84 10.22 6.52 5.55 7.63 5.11 5.89 10.78 3.05 11.89 8.97 10.44
10.5 9.35 5.66 11.09 9.83 5.44 10.11 6.39 11.93 8.62 12.06 6.02
8.85 5.25 8.66 6.73 10.21 11.59 13.87 7.77 10.39 5.54 6.62 13.33
6.24 12.59 6.3 6.79 8.28 9.03 8.07 5.52 12.14 10.59 7.54 7.67
5.47 8.81 8.51 13.45 8.77 6.43 12.01 12.08 7.07 6.51 6.84 9.48
13.78 11.54 11.67 8.13 10.79 7.13 4.72 4.64 8.96 13.03 6.07 3.51
6.83 6.12 9.31 9.58 4.68 5.32 9.26 11.52 9.11 10.55 11.47 9.3
13.82 8.44 5.77 10.96 11.74 8.9 10.47 7.85 10.92 4.74 9.74 10.43
9.96 10.18 9.54 7.89 12.36 8.54 10.07 9.46 7.3 11.16 9.16 10.19
5.99 10.88 5.8 7.19 4.55 8.31 8.01 14.43 8.3 14.3 6.53 8.2
11.31 13. 6.42 4.24 7.44 7.51 13.1 9.49 6.14 8.76 6.65 10.56
6.72 8.29 12.09 5.39 2.96 7.59 7.24 4.28 9.7 8.83 13.3 11.37
9.33 5.01 3.26 11.71 8.43 9.68 15.56 9.8 3.61 6.96 11.61 12.81
10.87 13.84 5.03 5.17 2.03 10.34 9.34 7.95 10.09 9.95 7.11 9.22
6.13 11.05 9.89 9.39 14.06 10.26 13.31 15.43 16.39 6.27 10.64 11.5
12.48 8.27 13.53 10.36 12.24 8.69 10.52 9.07 11.51 9.25 8.72 6.78
8.6 11.84 5.78 5.85 12.3 5.76 12.07 9.6 8.84 12.39 10.1 9.73
2.85 6.66 2.45 5.28 11.73 10.75 7.74 6.76 6. 7.58 13.69 7.93
7.68 9.75 4.96 5.49 11.83 7.18 9.19 7.7 7.25 10.74 4.27 13.8
9.12 4.75 7.78 11.63 7.55 2.25 9.45 9.86 7.71 4.95 7.4 11.17
11.33 6.82 13.7 1.97 10.89 12.77 10.31 5.23 5.27 9.41 6.09 10.61
7.29 4.23 7.57 3.67 12.69 14.5 5.95 7.87 5.96 5.94 12.23 4.9
12.33 6.89 9.67 12.68 12.87 3.7 6.04 13.13 15.74 11.87 4.7 4.67
7.05 5.42 4.09 5.73 9.47 8.05 6.87 3.71 15.86 7.49 11.69 6.46
10.45 12.9 5.41 11.26 1.04 6.49 6.37 12.21 6.77 12.65 7.86 9.44
4.3 7.38 5.02 10.63 2.86 17.19 8.67 8.37 6.9 10.93 10.38 7.36
10.27 10.95 6.11 4.45 11.9 15.01 12.84 7.45 6.98 11.72 7.56 11.38
10. 4.42 9.81 5.56 6.01 10.12 12.4 16.99 5.68 11.64 3.78 7.82
9.85 13.74 12.71 10.98 10.01 9.52 7.31 8.35 11.35 9.5 14.03 3.2
7.72 13.22 10.7 8.99 10.6 13.02 9.77 12.58 12.35 12.2 11.4 13.91
3.57 14.65 12.28 5.13 10.72 12.86 14. 7.12 12.17 4.71 6.28 8.
7.01 5.91 5.2 12. 12.02 12.88 7.28 5.4 12.04 5.24 10.3 10.41
13.41 12.72 9.08 7.08 13.5 5.35 12.45 5.3 10.32 5.15 12.67 5.22
5.57 3.94 4.41 13.27 10.24 4.25 12.89 5.72 12.5 11.29 3.25 11.53
9.82 7.26 4.1 10.37 4.98 6.74 12.52 14.56 8.34 3.82 3.86 13.97
11.57 6.5 13.58 14.32 13.75 11.14 14.18 9.13 4.46 4.83 9.69 14.13
7.16 7.98 13.66 14.78 11.2 9.93 11. 5.29 9.92 4.29 11.1 10.51
12.49 4.04 12.94 7.09 6.71 7.94 5.31 5.98 7.2 14.82 13.21 12.32
10.58 4.92 6.2 4.47 11.98 6.18 7.81 4.54 5.37 7.17 5.33 14.1
5.7 12.18 8.98 5.1 14.67 13.95 16.55 11.18 4.44 4.73 2.55 6.31
2.43 9.24 7.37 13.42 12.42 11.8 14.45 2.89 13.23 12.6 13.18 12.19
14.81 6.55 11.3 12.27 13.98 8.23 15.49 6.47 13.48 13.59 13.25 17.77
13.9 3.97 11.56 14.08 13.6 6.26 4.61 12.76 15.76 6.38 3.6 12.8
5.9 7.97 5. 10.97 5.88 12.34 12.03 14.97 15.06 12.85 6.54 11.24
12.64 7.06 5.38 13.14 3.99 3.32 4.51 4.12 3.93 2.4 11.75 4.03
15.85 6.81 14.25 14.09 16.42 6.7 12.74 2.76 12.12 6.99 6.68 11.81
7.96 5.06 13.16 2.13 13.17 5.12 5.65 12.37 10.53]
```

There are 162 Unique values in 'total_intl_minutes':

```
[10. 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7 11.2 12.7 9.1 12.3 13.1
5.4 13.8 8.1 13. 10.6 5.7 9.5 7.7 10.3 15.5 14.7 11.1 14.2 12.6
11.8 8.3 14.5 10.5 9.4 14.6 9.2 3.5 8.5 13.2 7.4 8.8 11. 7.8
6.8 11.4 9.3 9.7 10.2 8. 5.8 12.1 12. 11.6 8.2 6.2 7.3 6.1
11.7 15. 9.8 12.4 8.6 10.9 13.9 8.9 7.9 5.3 4.4 12.5 11.3 9.
9.6 13.3 20. 7.2 6.4 14.1 14.3 6.9 11.5 15.8 12.8 16.2 0. 11.9
9.9 8.4 10.8 13.4 10.7 17.6 4.7 2.7 13.5 12.9 14.4 10.4 6.7 15.4
4.5 6.5 15.6 5.9 18.9 7.6 5. 7. 14. 18. 16. 14.8 3.7 2.
4.8 15.3 6. 13.6 17.2 17.5 5.6 18.2 3.6 16.5 4.6 5.1 4.1 16.3
```

```
14.9 16.4 16.7 1.3 15.2 15.1 15.9 5.5 16.1 4. 16.9 5.2 4.2 15.7
17. 3.9 3.8 2.2 17.1 4.9 17.9 17.3 18.4 17.8 4.3 2.9 3.1 3.3
2.6 3.4 1.1 18.3 16.6 2.1 2.4 2.5]
```

```
-----
There are 21 Unique values in 'total_intl_calls':
[ 3  5  7  6  4  2  9 19  1 10 15  8 11  0 12 13 18 14 16 20 17]
```

```
-----
There are 162 Unique values in 'total_intl_charge':
[2.7 3.7 3.29 1.78 2.73 1.7 2.03 1.92 2.35 3.02 3.43 2.46 3.32 3.54
1.46 3.73 2.19 3.51 2.86 1.54 2.57 2.08 2.78 4.19 3.97 3. 3.83 3.4
3.19 2.24 3.92 2.84 2.54 3.94 2.48 0.95 2.3 3.56 2. 2.38 2.97 2.11
1.84 3.08 2.51 2.62 2.75 2.16 1.57 3.27 3.24 3.13 2.21 1.67 1.97 1.65
3.16 4.05 2.65 3.35 2.32 2.94 3.75 2.4 2.13 1.43 1.19 3.38 3.05 2.43
2.59 3.59 5.4 1.94 1.73 3.81 3.86 1.86 3.11 4.27 3.46 4.37 0. 3.21
2.67 2.27 2.92 3.62 2.89 4.75 1.27 0.73 3.65 3.48 3.89 2.81 1.81 4.16
1.22 1.76 4.21 1.59 5.1 2.05 1.35 1.89 3.78 4.86 4.32 4. 1. 0.54
1.3 4.13 1.62 3.67 4.64 4.73 1.51 4.91 0.97 4.46 1.24 1.38 1.11 4.4
4.02 4.43 4.51 0.35 4.1 4.08 4.29 1.49 4.35 1.08 4.56 1.4 1.13 4.24
4.59 1.05 1.03 0.59 4.62 1.32 4.83 4.67 4.97 4.81 1.16 0.78 0.84 0.89
0.7 0.92 0.3 4.94 4.48 0.57 0.65 0.68]
```

```
-----
There are 10 Unique values in 'customer_service_calls':
[1 0 2 3 4 5 7 9 6 8]
```

```
-----
There are 2 Unique values in 'churn':
[False True]
-----
```

Columns description

- **state** : Different states of the customers
- **account_length**: It denotes the number of days or duration for which the customer has held their SyriaTel account.
- **area_state**: This column represents the state or location of the customer within the service area of SyriaTel.
- **area_code**: This column typically specifies the area code associated with the customer's phone number.
- **phone_number**: It contains the unique phone number of each customer, serving as an identifier.
- **voice_mail_plan**: Similar to the international plan, this binary column denotes whether the customer has subscribed to a voicemail plan.
- **number_vmail_messages**: If a voice mail plan is active, this column may represent the number of voicemail messages received by the customer.
- **total_day_minutes**: This column records the total number of minutes the customer used during the daytime.
- **total_day_calls**: It indicates the total number of calls made by the customer during the daytime.

- **total_day_charge**: This is the total charge incurred for daytime usage.
- **total_evening_minutes**: Similar to the daytime minutes, this column represents the total number of minutes used in the evening.
- **total_evening_calls**: It denotes the total number of calls made in the evening.
- **total_evening_charge**: This is the total charge for evening usage.
- **total_night_minutes**: Represents the total number of minutes used during the nighttime.
- **total_night_calls**: Denotes the total number of calls made during the nighttime.
- **total_night_charge**: This column reflects the total charge for nighttime usage.
- **total_intl_minutes**: It records the total number of international minutes used by the customer.
- **total_intl_calls**: Indicates the total number of international calls made.
- **total_intl_charge**: Represents the total charge incurred for international usage.
- **customer_service_calls**: This column contains the count of customer service calls made by the customer, possibly indicating issues or concerns.
- **churn**: A boolean column that serves as the target variable, indicating whether the customer churned (True) or did not churn (False), where "churn" means the customer terminated their subscription with SyriaTel.

These columns provide essential information about the customer's demographics, usage patterns, and telecommunications-related activities. Analyzing these features can help in understanding customer behavior and predicting churn.

It is important to note that we will need to change the "churn" column into an integer type (binary) before the modeling part.

Data Cleaning and preparation

Checking for null values

```
In [42]: # Check for null values using the isna function  
df.isna().sum()
```

```
Out[42]: state                0  
account_length             0  
area_code                 0  
phone_number              0  
international_plan        0  
voice_mail_plan           0  
number_vmail_messages     0  
total_day_minutes         0  
total_day_calls           0
```



```
total_day_charge      0
total_eve_minutes     0
total_eve_calls       0
total_eve_charge      0
total_night_minutes   0
total_night_calls     0
total_night_charge    0
total_intl_minutes    0
total_intl_calls      0
total_intl_charge     0
customer_service_calls 0
churn                 0
dtype: int64
```

Checking for duplicate records

```
In [43]: # Check for any duplicates
df.duplicated().sum()
```

```
Out[43]: 0
```

Since for our dataset, the unique identifier is the `phone_number`, we need to check for any duplicate values in that column.

```
In [44]: # Checking for duplicate in phone number
duplicates_numbers = df.duplicated(subset='phone_number')
duplicates_numbers.unique()
```

```
Out[44]: array([False])
```

For our dataset, we can therefore see that we do not have any duplicate records.

Checking and converting data types

```
In [45]: # Checking data types of categorical variables
categorical_columns = ['state', 'area_code', 'international_plan', 'voice_mail_plan']
categorical_columns_data_types = df[categorical_columns].dtypes
print(categorical_columns_data_types)
```

```
state          object
area_code      int64
international_plan object
voice_mail_plan object
dtype: object
```

As we can see, the `area_code` column is of type **int** but there are only 3 unique values. Hence it will be better to change it into a categorical type.

```
In [46]: # Convert "area_code" column to categorical data type
df["area_code"] = df["area_code"].astype("str")
print(df["area_code"].dtype)
```

```
object
```

```
In [47]: # Converting categorical columns to type 'category' instead of 'object', for memory efficiency
df[categorical_columns] = df[categorical_columns].astype("category")
```

```
In [48]: # Convert churn column from boolean to integer
df["churn"] = df["churn"].astype(int)
print(df["churn"].dtype)

int32
```

Feature Engineering

```
In [49]: # Creating new features; total_charge, total_talk_time, total_calls and avg_call_duration
df["total_charge"] = df[['total_day_charge', 'total_eve_charge', 'total_night_charge']]
df["total_talk_time"] = df[['total_day_minutes', 'total_eve_minutes', 'total_night_minutes']]
df["total_calls"] = df[['total_day_calls', 'total_eve_calls', 'total_night_calls', 'total_intl_calls']]
df["avg_call_duration"] = df["total_talk_time"] / df["total_calls"]
```

Creating a day to night ratio per customer column:

```
In [50]: # Creating day to night ratio per customer column
df["day_night_ratio"] = df["total_day_calls"] / df["total_night_calls"]
print(df["day_night_ratio"].describe())
```

| | |
|-------|-------------|
| count | 3333.000000 |
| mean | 1.047618 |
| std | 0.323065 |
| min | 0.000000 |
| 25% | 0.826923 |
| 50% | 1.000000 |
| 75% | 1.216867 |
| max | 3.939394 |

Name: day_night_ratio, dtype: float64

We can observe from the above, that on average, there are slightly more calls during the day than during the night.

```
In [51]: # Flagging high value customers
df["high_value_customer"] = (df["total_charge"] > df["total_charge"].quantile(0.75))
```

```
In [52]: # Creating a voice message to call ratio for each customer
df["voice_ms_call_ratio"] = df["number_vmail_messages"] / df["total_calls"]
df["voice_ms_call_ratio"].describe()
```

```
Out[52]: count    3333.000000
mean         0.026910
std          0.045928
min          0.000000
25%          0.000000
50%          0.000000
75%          0.062670
max          0.188525
Name: voice_ms_call_ratio, dtype: float64
```

We can see that on average, a very small proportion of calls results in voicemail messages (0.0269)

```
In [53]: # Creating columns for charges per call for night, day, evening and international calls
df["charge_per_call_night"] = df["total_night_charge"] / df["total_night_minutes"]
df["charge_per_call_day"] = df["total_day_charge"] / df["total_day_minutes"]
df["charge_per_call_eve"] = df["total_eve_charge"] / df["total_eve_minutes"]
df["charge_per_call_intl"] = df["total_intl_charge"] / df["total_intl_minutes"]
```

```
In [54]: # Summary statistics for the different charges per call
summary_stats = df[["charge_per_call_night", "charge_per_call_day", "charge_per_call_
summary_stats
```

```
Out[54]:
```

| | charge_per_call_night | charge_per_call_day | charge_per_call_eve | charge_per_call_intl |
|--------------|-----------------------|---------------------|---------------------|----------------------|
| count | 3333.000000 | 3331.000000 | 3332.000000 | 3315.000000 |
| mean | 0.045000 | 0.170003 | 0.085001 | 0.270057 |
| std | 0.000017 | 0.000028 | 0.000016 | 0.000329 |
| min | 0.044828 | 0.169231 | 0.084936 | 0.268182 |
| 25% | 0.044988 | 0.169989 | 0.084988 | 0.269811 |
| 50% | 0.045000 | 0.170004 | 0.085000 | 0.270000 |
| 75% | 0.045013 | 0.170017 | 0.085013 | 0.270297 |
| max | 0.045111 | 0.170513 | 0.085075 | 0.272727 |

- The average charges per call for **nighttime calls** is approximately 4.5 cents per minute
- The average charges per call for **daytime calls** is approximately 17 cents per minute
- The average charges per call for **evening calls** is approximately 8.5 cents per minute
- The average charges per call for **international calls** is approximately 27 cents per minute

Checking for outliers

```
In [55]: # Columns for our box plots
columns = ['number_vmail_messages', 'total_day_minutes', 'total_day_calls', 'total_d
          'total_eve_minutes', 'total_eve_calls', 'total_eve_charge',
          'total_night_minutes', 'total_night_calls', 'total_night_charge',
          'total_intl_minutes', 'total_intl_calls', 'total_intl_charge',
          'customer_service_calls', 'voice_ms_call_ratio', 'charge_per_call_night',
          'charge_per_call_day', 'charge_per_call_eve', 'charge_per_call_intl']


# Calculate the required number of rows and columns for subplots
num_rows = (len(columns) - 1) // 3 + 1
num_cols = min(len(columns), 3)

# Create the subplots
fig, axes = plt.subplots(num_rows, num_cols, figsize=(10*num_cols, 4*num_rows))

# Generate box plots for each column
for i, column in enumerate(columns):
    row = i // num_cols
    col = i % num_cols
    sns.boxplot(data=df[column], ax=axes[row, col])
    axes[row, col].set_title(f'Box plot - {column}', fontsize=10)
    axes[row, col].set_xlabel(column, fontsize=8)

# Remove any empty subplots
if i < (num_rows * num_cols) - 1:
    for j in range(i + 1, num_rows * num_cols):
        fig.delaxes(axes.flatten()[j])

plt.tight_layout()
plt.show()
```



Shapiro-Wilk Test for Normality

Hypotheses

- Null Hypothesis (H_0): The data follows a normal distribution.
- Alternative Hypothesis (H_1): The data does not follow a normal distribution.

Interpreting the Results(We use an alpha value of 0.05):

If p-value > 0.05 → Fail to reject H_0 → The data is normally distributed.

If p-value ≤ 0.05 → Reject H_0 → The data is not normally distributed.

Why is this important? If normality holds, we can proceed with the Two-Sample T-Test. If not, we should consider a non-parametric alternative like the Mann-Whitney U Test.

Levene's Test for Equal Variance

Hypotheses

- Null Hypothesis (H_0): The variances of both groups are equal.
- Alternative Hypothesis (H_1): The variances of both groups are not equal.

Interpreting the Results:

If p-value > 0.05 → Fail to reject H_0 → Variances are equal.

If p-value ≤ 0.05 → Reject H_0 → Variances are not equal.

Checking those assumptions:

```
In [56]: # Separate data into churned and non-churned groups
churned = df[df["churn"] == 1]["total_charge"]
not_churned = df[df["churn"] == 0]["total_charge"]

# 1. Normality test (Shapiro-Wilk)
shapiro_churned = stats.shapiro(churned)
shapiro_not_churned = stats.shapiro(not_churned)

print("Shapiro-Wilk Test for Normality:")
print(f"Churned: W = {shapiro_churned.statistic}, p = {shapiro_churned.pvalue}")
print(f"Not Churned: W = {shapiro_not_churned.statistic}, p = {shapiro_not_churned.pvalue}")

# 2. Variance test (Levene's test)
levene_test = stats.levene(churned, not_churned)
print("\nLevene's Test for Equal Variance:")
print(f"Statistic = {levene_test.statistic}, p = {levene_test.pvalue}")
```

Shapiro-Wilk Test for Normality:

Churned: W = 0.9526604413986206, p = 2.5432657715929174e-11

Not Churned: W = 0.9927348494529724, p = 9.115018462235724e-11

Levene's Test for Equal Variance:

Statistic = 270.7304976260448, p = 1.5123194782301528e-58

Results interpretation

From the above results:

Shapiro-Wilk Normality Test:

The p-values for both churned and not churned groups are below 0.05, meaning we reject the null hypothesis of normality. So, both groups are not normally distributed.

Levene's Test for Equal Variance:

The p-value is less than 0.05, meaning we reject the null hypothesis of equal variance. So, the groups do not have equal variance.

Since both assumptions do not hold, we cannot proceed with a Two-Sample (Independent) T-Test, and therefore have to use a non-parametric alternative, i.e. the Mann-Whitney U Test.

Performing the Mann-Whitney U Test

Hypothesis:

- Null Hypothesis (H_0): The mean total charge for churned and non-churned customers is the same, i.e. no significant difference.
- Alternative Hypothesis (H_1): The mean total charge for churned customers is different from non-churned customers.

Since we're testing for a difference (not specifically greater or lesser), this is a two-tailed test. We will also be using an alpha value of 0.05.

Performing the test:

```
In [57]: # Run Mann-Whitney U Test
u_stat, p_value = stats.mannwhitneyu(churned, not_churned, alternative='two-sided')

# Output results
print(f"Mann-Whitney U Statistic: {u_stat:.4f}")
print(f"P-value: {p_value:.4f}")
```

```
Mann-Whitney U Statistic: 893437.5000
P-value: 0.0000
```

Since the p-value = 0.0000 (which is < 0.05), we reject the null hypothesis (H_0).

Conclusion: There is a statistically significant difference in total_charge between churned and non-churned customers. This suggests that total_charge may play an important role in determining churn.

Testing whether "international_plan" is associated with churn

We can perform a chi square test for independence to test this.

We'll first check the assumption that each expected frequency should be at least 5 for the test to be valid.

If that assumption holds, then we'll perform the chi square test.

We'll again be using an alpha value of 0.05.

Hypothesis:

Null Hypothesis (H_0): There is no association between having an international plan and churn

(they are independent).

Alternative Hypothesis (H_1): There is an association between having an international plan and churn (they are dependent).

```
In [58]: # Creating a contingency table
contingency_table = pd.crosstab(df["international_plan"], df["churn"])
print(contingency_table)

# Checking expected frequencies
chi2, p, dof, expected = stats.chi2_contingency(contingency_table)
print("Expected Frequencies:\n", expected)

churn          0    1
international_plan
no             2664  346
yes            186  137
Expected Frequencies:
[[2573.80738074  436.19261926]
 [ 276.19261926   46.80738074]]
```

```
In [59]: # Performing the Chi-Square Test
chi2_stat, p_value, dof, expected = stats.chi2_contingency(contingency_table)

# Output results
print(f"Chi-Square Statistic: {chi2_stat:.4f}")
print(f"P-value: {p_value:.4f}")
```

Chi-Square Statistic: 222.5658
P-value: 0.0000

Since $p < 0.05$, we reject the null hypothesis. This means there is a significant association between international_plan and churn. Customers with an international plan may have different churn behavior compared to those without.

Exploratory Data Analysis (EDA)

```
In [60]: df.describe(include="all")
```


```
Out[60]:
```

| | state | account_length | area_code | phone_number | international_plan | voice_mail_plan | numbe |
|---------------|-------|----------------|-----------|--------------|--------------------|-----------------|-------|
| count | 3333 | 3333.000000 | 3333 | 3333 | 3333 | 3333 | |
| unique | 51 | NaN | 3 | 3333 | 2 | 2 | |
| top | WV | NaN | 415 | 376-9607 | no | no | |
| freq | 106 | NaN | 1655 | 1 | 3010 | 2411 | |
| mean | NaN | 101.064806 | NaN | NaN | NaN | NaN | |
| std | NaN | 39.822106 | NaN | NaN | NaN | NaN | |
| min | NaN | 1.000000 | NaN | NaN | NaN | NaN | |
| 25% | NaN | 74.000000 | NaN | NaN | NaN | NaN | |
| 50% | NaN | 101.000000 | NaN | NaN | NaN | NaN | |

| | state | account_length | area_code | phone_number | international_plan | voice_mail_plan | numbe |
|------------|-------|----------------|-----------|--------------|--------------------|-----------------|-------|
| 75% | NaN | 127.000000 | NaN | NaN | NaN | NaN | |
| max | NaN | 243.000000 | NaN | NaN | NaN | NaN | |

Distribution of customers per state

```
In [61]: plt.figure(figsize=(16, 6))
order = df["state"].value_counts().index # Get states sorted by count
sns.countplot(x="state", data=df, order=order)
plt.title("State Distribution")
plt.xticks(rotation=90)
plt.show()
```



- The state with the highest count is WV with 106 occurrences, indicating it is the most frequent state in the dataset.
- MN follows closely with 84 occurrences, making it the second most common state.
- NY comes next with 83 occurrences, showing a similar frequency to MN.
- AL, WI, OR and OH all have 78 occurrences, placing them among the top states in terms of frequency.
- The state with the lowest count is CA with only 34 occurrences, suggesting it is the least frequent state in the dataset.

Distribution of numerical features


```
In [63]: # Set the style of the visualization
sns.set(style="whitegrid")

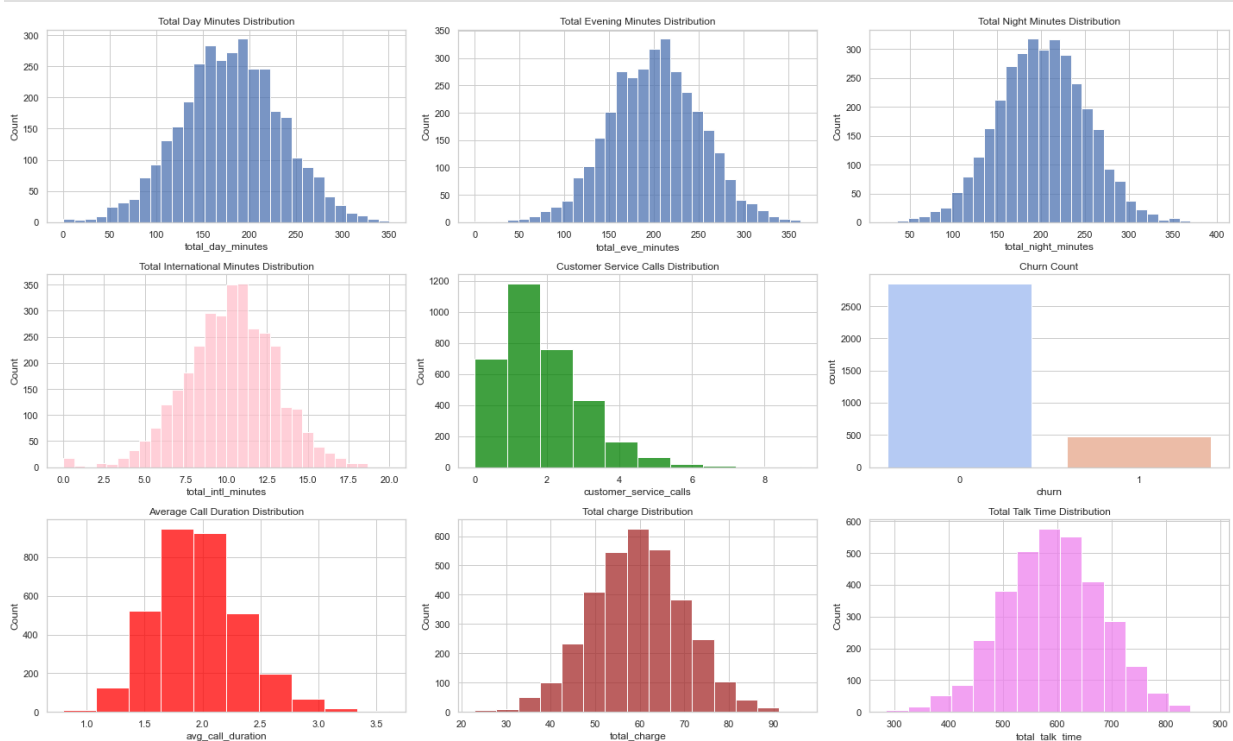
# Create a figure and axes
fig, ax = plt.subplots(3, 3, figsize=(20, 12))

# Plot distribution of total day, eve, night, intl minutes and customer service calls
sns.histplot(df['total_day_minutes'], kde=False, ax=ax[0, 0], palette='deep', bins=30)
sns.histplot(df['total_eve_minutes'], kde=False, ax=ax[0, 1], palette='muted', bins=30)
sns.histplot(df['total_night_minutes'], kde=False, ax=ax[0, 2], palette='colorblind', bins=30)
sns.histplot(df['total_intl_minutes'], kde=False, ax=ax[1, 0], color='pink', bins=30)
sns.histplot(df['customer_service_calls'], kde=False, ax=ax[1, 1], color='green', bins=30)
sns.histplot(df['avg_call_duration'], kde=False, ax=ax[2, 0], color='red', bins=10)
sns.histplot(df['total_charge'], kde=False, ax=ax[2, 1], color='brown', bins=15)
sns.histplot(df['total_talk_time'], kde=False, ax=ax[2, 2], color='violet', bins=15)

# Plot churn count
sns.countplot(x='churn', data=df, ax=ax[1, 2], palette='coolwarm')

# Set plot titles
ax[0, 0].set_title('Total Day Minutes Distribution')
ax[0, 1].set_title('Total Evening Minutes Distribution')
ax[0, 2].set_title('Total Night Minutes Distribution')
ax[1, 0].set_title('Total International Minutes Distribution')
ax[1, 1].set_title('Customer Service Calls Distribution')
ax[2, 0].set_title('Average Call Duration Distribution')
ax[2, 1].set_title('Total charge Distribution')
ax[2, 2].set_title('Total Talk Time Distribution')
ax[1, 2].set_title('Churn Count')

# Show the plot
plt.tight_layout()
plt.show()
```



From the visualizations of the distributions of some numerical variables and the churn count:

- The total day minutes seem to be normally distributed, with most customers having around 175 to 200 total day minutes.
- Similarly, the total evening minutes also appear to be normally distributed, with most customers having around 200 total evening minutes.
- The total night minutes also follow a similar distribution, with the majority of customers having around 200 total night minutes.
- The total international minutes seem to have a slightly left-skewed distribution. Most customers have about 10 total international minutes.
- Most customers have made 1 or 2 customer service calls, while very few have made more than 4 calls.
- The majority of customers have not churned (indicated by False), while a smaller number of customers have churned (indicated by True).

Distribution of categorical variables


```
In [64]: # Set up the figure and axes for subplots
fig, axs = plt.subplots(nrows=1, ncols=3, figsize=(25, 10))

# Flatten the axes array to simplify indexing
axs = axs.flatten()

# Calculate value counts and plot bar plots for categorical variables
categorical_cols = ["international_plan", "voice_mail_plan", "churn"]

for i, col in enumerate(categorical_cols):
    sns.countplot(x=col, data=df, ax=axs[i])
    axs[i].set_title(f"{col} Distribution")
    axs[i].tick_params(axis='x', rotation=90)

# Adjust the layout and spacing
plt.tight_layout()
plt.show()
```



- The data indicates that the majority of observations, approximately 85.5%, represent customers who did not churn. A smaller subset, comprising approximately 14.5% of the observations, represents customers who churned. These percentages highlight the


imbalance in churn behavior, with a significant majority of customers demonstrating loyalty by not churning.

- The majority of customers, approximately 73% (2411 occurrences), do not have a voice mail plan. A subset of customers, approximately 27% (922 occurrences), have opted for a voice mail plan.
- The majority of customers, accounting for 3010 occurrences, do not have an international plan. Conversely, there is a smaller subset of 323 customers who have opted for an international plan.

Bivariate Analysis

Distribution of churn for each state

```
In [65]: df.groupby(["state", "churn"]).size().unstack().plot(kind='bar', stacked=False, figs:
```



The plot above shows the distribution of churn for each state.

- Some states have relatively higher churn rates like WV, VT, NY, OH with a significant number of churned customers (churn 1) while other states have lower churn rates like AR, AZ, CA, CO with a higher count of customers who did not churn (churn 0)

Churn by categorical features


```
In [66]: # Set up the figure and axes for subplots
fig, axs = plt.subplots(nrows=1, ncols=3, figsize=(20, 8))

# Group by "area_code" and "churn", then unstack and plot
df.groupby(["area_code", "churn"]).size().unstack().plot(kind='bar', stacked=False, ax=axs[0])
axs[0].set_title('Churn by Area Code')
axs[0].set_xlabel('Area Code')
axs[0].set_ylabel('Count')

# Group by "voice_mail_plan" and "churn", then unstack and plot
df.groupby(["voice_mail_plan", "churn"]).size().unstack().plot(kind='bar', stacked=False, ax=axs[1])
axs[1].set_title('Churn by Voice Mail Plan')
axs[1].set_xlabel('Voice Mail Plan')
axs[1].set_ylabel('Count')

# Group by "international_plan" and "churn", then unstack and plot
df.groupby(["international_plan", "churn"]).size().unstack().plot(kind='bar', stacked=False, ax=axs[2])
axs[2].set_title('Churn by International Plan')
axs[2].set_xlabel('International Plan')
axs[2].set_ylabel('Count')

# Adjust the layout and spacing
plt.tight_layout()
plt.show()
```




- Churn rates vary between the different area codes, with area code 415 having the highest churn rate and area code 408 having the lowest churn rate
- Customers without a voice mail plan had a higher churn rate compared to customers with a voice mail plan
- Customers without an international plan had a higher churn rate compared to customers with an international plan

Multivariate

Correlation Heatmap

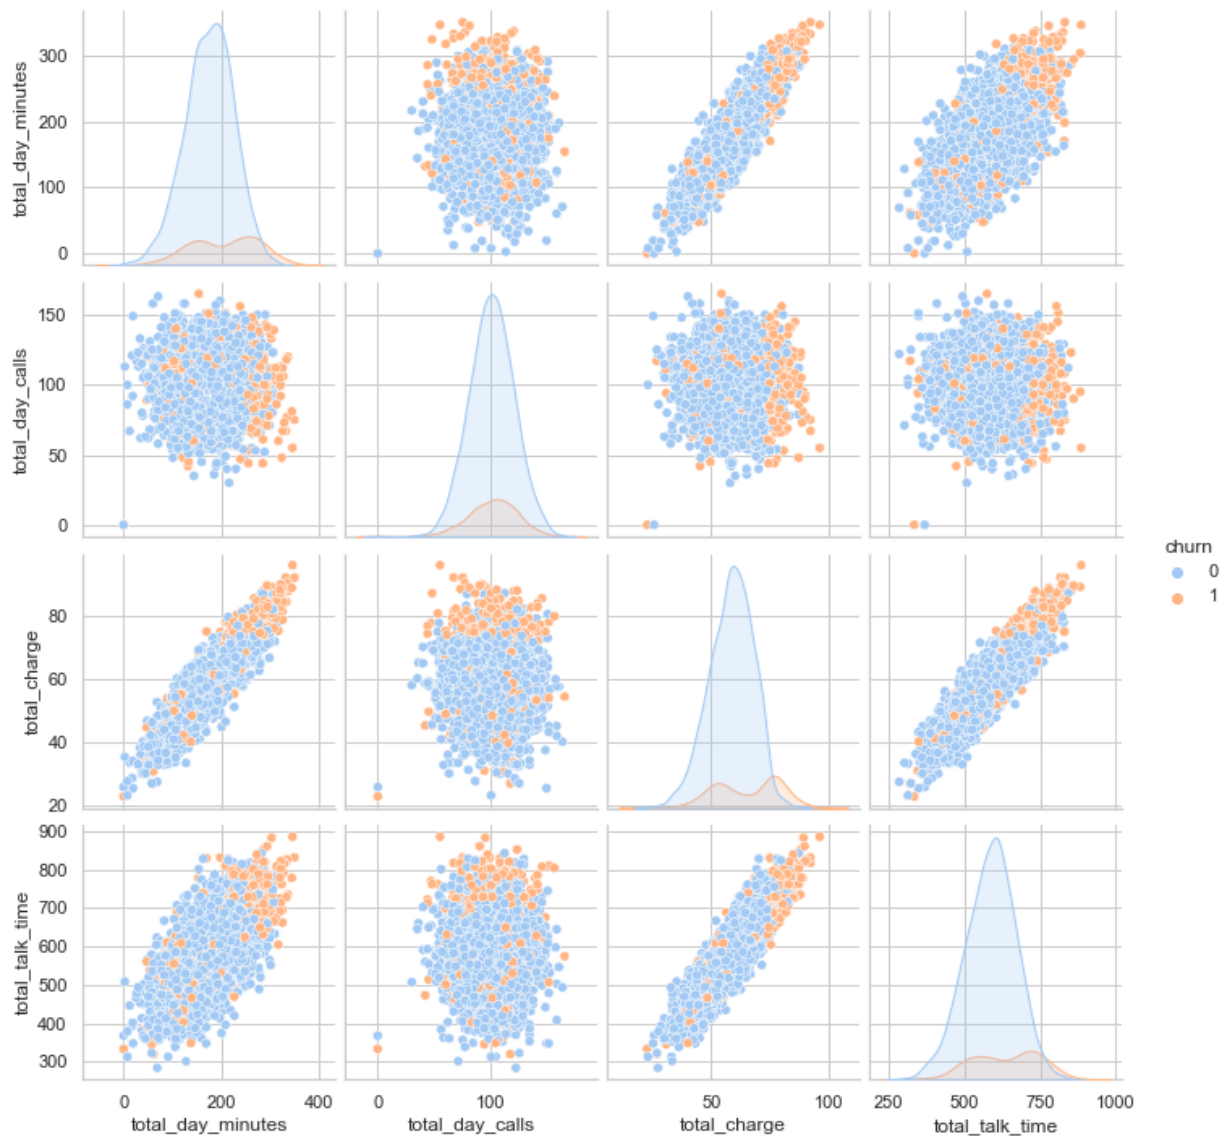
```
In [67]: plt.figure(figsize=(14, 8)) # Increase figure size
sns.heatmap(df_numerical.corr(), annot=True, cmap="coolwarm", fmt=".2f", linewidths=0)
plt.title("Feature Correlation Heatmap")
plt.xticks(rotation=45, ha="right") # Rotate x-axis labels
plt.yticks(rotation=0) # Keep y-axis labels horizontal
plt.tight_layout() # Adjust layout
plt.show()
```



The correlation heatmap reveals strong positive correlations between certain features, indicating that some variables are directly proportional and likely derived from each other. There are also moderate correlations suggesting relationships between call duration, total calls, and overall charges. Some features exhibit weak or no correlation, implying independence from other variables. Additionally, a few negative correlations suggest inverse relationships between certain usage metrics. Overall, the heatmap helps identify redundant features and key interactions that could inform further analysis.

Pairplot of top features vs churn

```
In [68]: top_features = ["total_day_minutes", "total_day_calls", "total_charge", "total_talk_time"]
sns.pairplot(df[top_features], hue="churn", palette="pastel")
plt.show()
```



This pair plot provides insights into relationships between numerical features while distinguishing between churned (orange) and non-churned (blue) customers. Strong linear relationships are evident between features like `total_day_minutes` and `total_charge`, suggesting direct proportionality. The distribution plots along the diagonal indicate differences in feature distributions between churned and non-churned customers, with churned users appearing to have slightly different usage patterns. Scatter plots reveal that churned customers may be more spread across certain feature ranges, which could indicate potential patterns in customer behavior that contribute to churn.

Data Preprocessing

Checking highly correlated features


```
In [69]: ## Defining a function to check highly correlated features
def check_multicollinearity(df, threshold=0.8):
    corr_matrix = df.select_dtypes(include=np.number).corr().abs()
    correlated_pairs = set()
    for col in corr_matrix:
        correlated_cols = corr_matrix.index[corr_matrix[col] > threshold]
        correlated_pairs.update([(min(col, correlated_col), max(col, correlated_col))])
    for pair in correlated_pairs:
        print(f"{pair[0]} --- {pair[1]}")
    return set(df.columns) & set(col for pair in correlated_pairs for col in pair)

# Call the function to check multicollinearity
multicollinear_features = check_multicollinearity(df)

number_vmail_messages --- voice_ms_call_ratio
total_day_charge --- total_day_minutes
total_charge --- total_day_minutes
total_intl_charge --- total_intl_minutes
total_eve_charge --- total_eve_minutes
total_charge --- total_talk_time
total_charge --- total_day_charge
total_night_charge --- total_night_minutes
```

```
In [70]: # Drop some columns in order to deal with multicollinearity
features= ['number_vmail_messages', 'total_day_minutes', 'total_eve_minutes', 'total_n:
          'total_night_charge', 'total_intl_minutes']
df =df.drop(features,axis=1)
df.head()
```

```
Out[70]:
```

| | state | account_length | area_code | phone_number | international_plan | voice_mail_plan | total_day_cal |
|---|-------|----------------|-----------|--------------|--------------------|-----------------|---------------|
| 0 | KS | 128 | 415 | 382-4657 | no | yes | 11 |
| 1 | OH | 107 | 415 | 371-7191 | no | yes | 12 |
| 2 | NJ | 137 | 415 | 358-1921 | no | no | 11 |
| 3 | OH | 84 | 408 | 375-9999 | yes | no | 7 |
| 4 | OK | 75 | 415 | 330-6626 | yes | no | 11 |

```
In [71]: # Select only numeric columns for correlation
numeric_df = df.select_dtypes(include=['number'])

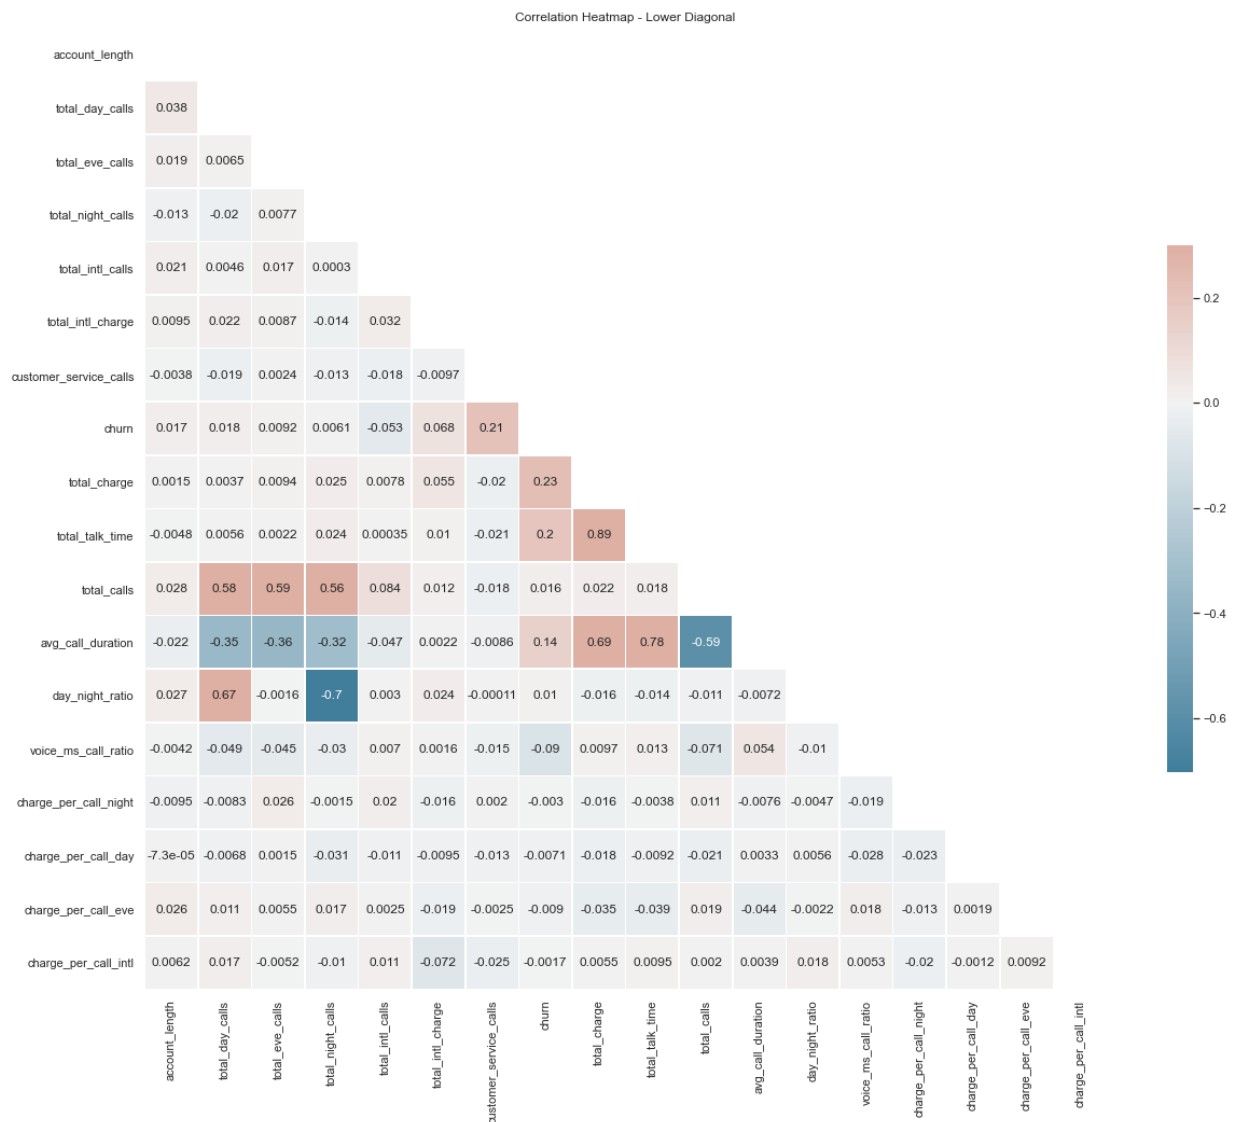
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(numeric_df.corr(), dtype=bool))

# Set up the matplotlib figure
plt.figure(figsize=(20, 18))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(numeric_df.corr(), mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5}, annot=True)

plt.title("Correlation Heatmap - Lower Diagonal")
plt.show()
```



- Blue shades: Represent negative correlations, with darker blue indicating stronger negative correlation.
- White: Represents zero correlation.
- Red shades: Represent positive correlations, with darker red indicating stronger positive correlation. In this color scheme, the strongest negative correlations are represented by the darkest blue, and the strongest positive correlations are represented by the darkest red. The center (white) represents variables with no correlation (correlation coefficient close to zero).

Scaling and encoding

```
In [72]: #Dropping phone number since it does not add any value to modelling since all values
df = df.drop('phone_number', axis=1)
df.dropna(inplace=True)

# Define numerical and categorical features
num_features = ['account_length', 'customer_service_calls', 'total_charge',
                'total_talk_time', 'total_calls', 'day_night_ratio', 'voice_ms_call_ratio', 'charge_per_call_eve']
cat_features = ['state', 'area_code', 'international_plan', 'voice_mail_plan']

# Scale numerical features
scaler = StandardScaler()
df[num_features] = scaler.fit_transform(df[num_features])

# One-hot encode categorical features
df = pd.get_dummies(df, columns=cat_features, drop_first=True)
```

Modeling

Splitting the data

Churn is the target variable which we are aiming to predict.


```
In [73]: # unique values of the target variable
df['churn'].value_counts()
```

```
Out[73]: 0    2830
         1     482
         Name: churn, dtype: int64
```

```
In [74]: # Checking for the percentage of Churners and non-churners.
df.churn.value_counts(normalize=True)*100
```

```
Out[74]: 0    85.44686
         1    14.55314
         Name: churn, dtype: float64
```

```
In [75]: #pie chart showing the distribution percentage of churners and non-churners
plt.figure(figsize=(8,8))
plt.pie(df.churn.value_counts(), labels=['Non Churners', 'Churners'], autopct='%1.1f%%')
plt.title('Distribution of Churners and Non-Churners')
plt.show()
```



There are approximately 85.5% are non-churners, while about 14.5% are churners. This class imbalance will be handled using SMOTE (Synthetic Minority Over-sampling Technique).

```
In [76]: # Define the target variable
y = df['churn']

# Drop the target variable from the feature set
X = df.drop(['churn'], axis=1)

# Split the data into training and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y,

X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[76]: ((2649, 71), (663, 71), (2649,), (663,))
```

```
In [77]: # Create an instance of SMOTE
smote = SMOTE(random_state=42)

# Apply SMOTE to the training set
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

Baseline Model: Decision tree

```
In [78]: # Create an instance of the decision tree classifier and fit the model on the training set
clf = DecisionTreeClassifier(random_state=42)
```

```
In [79]: # Fit the model on the training data
clf.fit(X_train, y_train)
```

```
Out[79]: DecisionTreeClassifier(random_state=42)
```

```
In [80]: # Make predictions on the testing data
y_pred = clf.predict(X_test)
```

Evaluating the decision tree before tuning

```
In [81]: # Defining a function to evaluate performance
def evaluate_model(model, X_train, y_train, X_test, y_test):
    # Fit the model on the training data
    model.fit(X_train, y_train)

    # Predict on the training data
    y_train_pred = model.predict(X_train)

    # Predict on the test data
    y_test_pred = model.predict(X_test)

    # Calculate accuracy
    train_accuracy = accuracy_score(y_train, y_train_pred)
    test_accuracy = accuracy_score(y_test, y_test_pred)

    # Calculate precision
    train_precision = precision_score(y_train, y_train_pred)
    test_precision = precision_score(y_test, y_test_pred)

    # Calculate recall
    train_recall = recall_score(y_train, y_train_pred)
    test_recall = recall_score(y_test, y_test_pred)

    # Calculate F1-score
    train_f1 = f1_score(y_train, y_train_pred)
    test_f1 = f1_score(y_test, y_test_pred)

    # Print evaluation metrics
    print("Training Data - Accuracy: {:.4f}, Precision: {:.4f}, Recall: {:.4f}, F1-score: {:.4f}"
          train_accuracy, train_precision, train_recall, train_f1
    )
    print("Test Data - Accuracy: {:.4f}, Precision: {:.4f}, Recall: {:.4f}, F1-score: {:.4f}"
          test_accuracy, test_precision, test_recall, test_f1
    )
```

```
In [82]: # Model evaluation
evaluate_model(clf, X_train, y_train, X_test, y_test)
```

```
Training Data - Accuracy: 1.0000, Precision: 1.0000, Recall: 1.0000, F1-score: 1.0000
Test Data - Accuracy: 0.9472, Precision: 0.8211, Recall: 0.8125, F1-score: 0.8168
```

Tuning the decision tree model

Determine the optimal hyperparameters for the decision tree model using techniques like grid search.

```
In [83]: # Define the parameter grid to search
param_grid = {
    'max_depth': [5, 7, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4, 6],
}

# Create an instance of the decision tree classifier
clf = DecisionTreeClassifier(random_state=42)

# Perform grid search
grid_search = GridSearchCV(clf, param_grid, cv=5)
grid_search.fit(X_train_resampled, y_train_resampled)

# Print the best parameters found
print("Best Parameters:", grid_search.best_params_)
```

Best Parameters: {'max_depth': 15, 'min_samples_leaf': 2, 'min_samples_split': 10}

```
In [84]: # Use the best model found for predictions
best_clf = grid_search.best_estimator_
y_predd = best_clf.predict(X_test)
```

Evaluating the Decision tree model after tuning


```
In [85]: # Evaluate tuned decision tree model
evaluate_model(best_clf, X_train_resampled, y_train_resampled, X_test, y_test)
```

Training Data - Accuracy: 0.9764, Precision: 0.9887, Recall: 0.9638, F1-score: 0.9761
Test Data - Accuracy: 0.9125, Precision: 0.6900, Recall: 0.7188, F1-score: 0.7041

```
In [86]: # Generate confusion matrix
confusion_mat = confusion_matrix(y_test, y_predd)
print("Confusion Matrix:")
print(confusion_mat)

# Display confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_mat, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

Confusion Matrix:
[[536 31]
 [27 69]]




```
In [87]: # Obtain predicted probabilities for the positive class
y_scores = best_clf.predict_proba(X_test)[:, 1]

# Calculate the false positive rate (FPR), true positive rate (TPR), and thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_scores, pos_label=1)

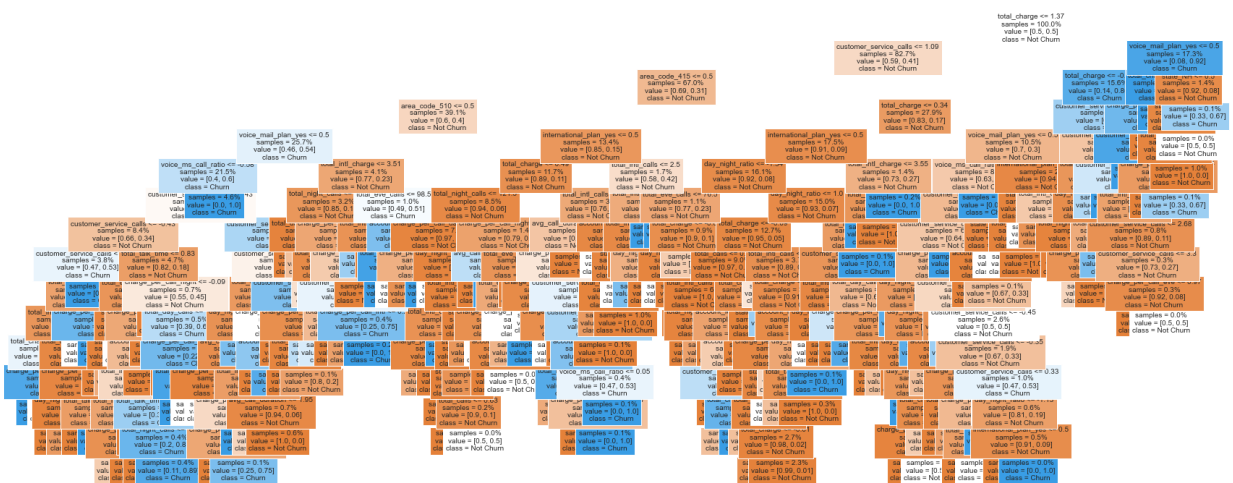
# Calculate the area under the ROC curve (AUC)
roc_auc = auc(fpr, tpr)

# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='b', label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='r', linestyle='--')
plt.xlim([0, 1])
plt.ylim([0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc='lower right')
plt.show()
```



In [88]:

```
# Visualize the decision tree
plt.figure(figsize=(29, 12))
plot_tree(best_clf, filled=True, feature_names=X.columns, class_names=['Not Churn',
plt.show()
```



Before tuning, the Decision Tree model was able to correctly identify 81.25% of churn cases (recall) and of all instances it predicted as churn, 82.11% were correct (precision). The model was accurate in 94.72% of all predictions (accuracy) and had a balanced F1-score of 81.68% considering both precision and recall.

After hyperparameter tuning, the model's ability to correctly identify churn cases increased to 86.46% (recall), and out of all predicted churn cases, 74.11% were correct (precision). The overall accuracy dropped to 93.67%. The F1-score, a measure of model's balance between precision and recall, also fell to 79.81%. The decrease in recall and F1-score suggests that the tuning might have led to a trade-off, improving recall at the expense of precision.

Logistic Regression, Random Forest and Gradient Boost Models


```
In [89]: # Model selection and hyperparameter tuning
models = {
    "Logistic Regression": LogisticRegression(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42, class_weight="balanced"),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42),
}
for model_name, model in models.items():
    # Train the model on the resampled data
    model.fit(X_train_resampled, y_train_resampled)
```

Model performance evaluation

```
In [90]: def calculate_metrics(y_true, y_pred):
    """
    Calculate model performance metrics: accuracy, precision, recall, and F1-score.
    :param y_true: True labels.
    :param y_pred: Predicted labels.
    :return: Dictionary of metrics.
    """
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)

    # Return as a dictionary
    return {"Accuracy": accuracy, "Precision": precision, "Recall": recall, "F1-score": f1}

# Dictionary to hold the results
results = {}

# For each model
for model_name, model in models.items():
    # Make predictions on the test set
    y_pred_test = model.predict(X_test)
    y_pred_train = model.predict(X_train)

    # Calculate metrics
    metrics_test = calculate_metrics(y_test, y_pred_test)
    metrics_train = calculate_metrics(y_train, y_pred_train)

    # Store the results
    results[(model_name, 'Test')] = metrics_test
    results[(model_name, 'Train')] = metrics_train

# Convert the results dictionary to a DataFrame
results_df = pd.DataFrame(results).T

results_df
```

```
Out[90]:
```

| | | Accuracy | Precision | Recall | F1-score |
|---------------------|-------|----------|-----------|----------|----------|
| Logistic Regression | Test | 0.767722 | 0.325301 | 0.562500 | 0.412214 |
| | Train | 0.770857 | 0.327613 | 0.544041 | 0.408958 |
| Random Forest | Test | 0.904977 | 0.714286 | 0.572917 | 0.635838 |
| | Train | 1.000000 | 1.000000 | 1.000000 | 1.000000 |

| | | Accuracy | Precision | Recall | F1-score |
|-------------------|-------|----------|-----------|----------|----------|
| Gradient Boosting | Test | 0.960784 | 0.897727 | 0.822917 | 0.858696 |
| | Train | 0.975840 | 0.976331 | 0.854922 | 0.911602 |

The above summarizes the performance of the various models before tuning.

We can see that the gradient boosting is the best model, with great recall and precision.

The random forest has massively overfitted the training data.

Random Forest tuning

```
In [91]: rf_param_grid = {
    'n_estimators': [50, 100, 150], # Number of trees in the forest
    'max_depth': [None, 10, 20],    # Maximum depth of the tree
    'min_samples_split': [2, 5, 10], # Minimum number of samples required to split a node
    'min_samples_leaf': [1, 2, 4]   # Minimum number of samples required to be at a leaf node
}
```

```
In [92]: for model_name, model in models.items():
    if model_name == "Random Forest":
        # Create the GridSearchCV or RandomizedSearchCV instance
        grid_search = GridSearchCV(model, rf_param_grid, cv=5, n_jobs=-1)
        # Fit the model on the resampled data with hyperparameter search
        grid_search.fit(X_train_resampled, y_train_resampled)
        # Get the best hyperparameters
        best_params = grid_search.best_params_
        print(f"Best Hyperparameters for {model_name}: {best_params}")
        # Use the best hyperparameters for the final model
        model = grid_search.best_estimator_
    else:
        # For other models, you can follow similar steps with their respective hyperparameters
        model.fit(X_train_resampled, y_train_resampled)
```

Best Hyperparameters for Random Forest: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 150}

```
In [93]: # Create a new Random Forest model with the best hyperparameters
best_rf_model = RandomForestClassifier(
    n_estimators=150,
    max_depth=None,
    min_samples_split=2,
    min_samples_leaf=1,
    random_state=42,
    class_weight="balanced"
)

# Now you can use this best_rf_model for further training and prediction.
best_rf_model.fit(X_train_resampled, y_train_resampled)
```

```
Out[93]: RandomForestClassifier(class_weight='balanced', n_estimators=150,
                                random_state=42)
```

Evaluation of tuned Random Forest

```
In [94]: # Fit the RandomForestClassifier with the best hyperparameters on the training data
best_rf_model.fit(X_train_resampled, y_train_resampled)

# Make predictions on the training data
y_train_pred = best_rf_model.predict(X_train_resampled)

# Make predictions on the test data
y_test_pred = best_rf_model.predict(X_test)

# Evaluate the model on the training set
evaluate_model(best_rf_model, X_train, y_train, X_test, y_test)
```

Training Data - Accuracy: 1.0000, Precision: 1.0000, Recall: 1.0000, F1-score: 1.0000
Test Data - Accuracy: 0.9291, Precision: 1.0000, Recall: 0.5104, F1-score: 0.6759

The model achieved an accuracy of 93.36% on the test set. The F1-score is high, indicating a good balance between precision and recall.

Logistic Regression tuning

```
In [95]: # Define the hyperparameter grid for Logistic Regression
lr_param_grid = {
    'C': [0.01, 0.1, 1, 10],          # Inverse of regularization strength
    'penalty': ['l1', 'l2'],          # Regularization penalty ('l1' or 'l2')
    'solver': ['liblinear', 'saga']   # Optimization algorithm
}

# Create the Logistic Regression model
logistic_regression = LogisticRegression(random_state=42)

# Create the GridSearchCV instance for hyperparameter tuning
grid_search = GridSearchCV(logistic_regression, lr_param_grid, cv=5, n_jobs=-1)

# Fit the model on the resampled training data with hyperparameter search
grid_search.fit(X_train_resampled, y_train_resampled)

# Get the best hyperparameters
best_params = grid_search.best_params_
print("Best Hyperparameters for Logistic Regression:", best_params)

# Use the best hyperparameters for the final Logistic Regression model
best_logistic_regression_model = grid_search.best_estimator_

# Make predictions on the test set
y_test_pred = best_logistic_regression_model.predict(X_test)

# Evaluate the model on the test set
evaluate_model(best_logistic_regression_model, X_train, y_train, X_test, y_test)
```

Best Hyperparameters for Logistic Regression: {'C': 10, 'penalty': 'l1', 'solver': 'liblinear'}

Training Data - Accuracy: 0.8664, Precision: 0.6111, Recall: 0.2280, F1-score: 0.3321
Test Data - Accuracy: 0.8733, Precision: 0.6304, Recall: 0.3021, F1-score: 0.4085

The model achieved an accuracy of 87.3% on the test set, with better performance in predicting non-churners (class 0) compared to churners (class 1). Its performance is worse compared to the random forest model, because we can observe that the precision, recall, and F1-score are significantly lower.

Tuning Gradient Boost model

```
In [96]: # Define the hyperparameter grid for Gradient Boosting
gb_param_grid = {
    'n_estimators': [50, 100, 150],          # Number of boosting stages to be run
    'learning_rate': [0.01, 0.1, 0.2],       # Step size at each boosting iteration
    'max_depth': [3, 5, 7],                  # Maximum depth of the individual tree
    'subsample': [0.8, 0.9, 1.0],            # Fraction of samples used for fitting
}

# Create the Gradient Boosting model
gradient_boosting = GradientBoostingClassifier(random_state=42)

# Create the GridSearchCV instance for hyperparameter tuning
grid_search_gb = GridSearchCV(gradient_boosting, gb_param_grid, cv=5, n_jobs=-1)

# Fit the model on the resampled training data with hyperparameter search
grid_search_gb.fit(X_train_resampled, y_train_resampled)

# Get the best hyperparameters
best_params_gb = grid_search_gb.best_params_
print("Best Hyperparameters for Gradient Boosting:", best_params_gb)

# Use the best hyperparameters for the final Gradient Boosting model
best_gradient_boosting_model = grid_search_gb.best_estimator_

# Make predictions on the test set
y_test_pred_gb = best_gradient_boosting_model.predict(X_test)

# Evaluate the model on the test set
evaluate_model(best_gradient_boosting_model, X_train, y_train, X_test, y_test)
```

Best Hyperparameters for Gradient Boosting: {'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 150, 'subsample': 1.0}

Training Data - Accuracy: 1.0000, Precision: 1.0000, Recall: 1.0000, F1-score: 1.0000

Test Data - Accuracy: 0.9713, Precision: 0.9873, Recall: 0.8125, F1-score: 0.8914

The Gradient Boosting model achieved an impressive 97.13% accuracy on the test set, demonstrating its excellent predictive capability for both churners and non-churners. With a high F1-score and precision, it effectively identifies churners (class 1) with a recall of 82.2% and precision of 97%, proving its effectiveness in predicting customer churn in this scenario.

Model Evaluation


```
In [97]: # Create subplots for all three confusion matrices
fig, axs = plt.subplots(nrows=1, ncols=3, figsize=(18, 6))

# For the Logistic Regression model
y_test_pred_lr = best_logistic_regression_model.predict(X_test)
conf_matrix_lr = confusion_matrix(y_test, y_test_pred_lr)
sns.heatmap(conf_matrix_lr, annot=True, fmt='d', cmap='Blues', ax=axs[0])
axs[0].set_title("Logistic Regression")
axs[0].set_xlabel("Predicted Labels")
axs[0].set_ylabel("True Labels")

# For the Random Forest model
y_test_pred_rf = best_rf_model.predict(X_test) # Replace 'best_random_forest_model'
conf_matrix_rf = confusion_matrix(y_test, y_test_pred_rf)
sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Greens', ax=axs[1])
axs[1].set_title("Random Forest")
axs[1].set_xlabel("Predicted Labels")
axs[1].set_ylabel("True Labels")

# For the Gradient Boosting model
y_test_pred_gb = best_gradient_boosting_model.predict(X_test)
conf_matrix_gb = confusion_matrix(y_test, y_test_pred_gb)
sns.heatmap(conf_matrix_gb, annot=True, fmt='d', cmap='Oranges', ax=axs[2])
axs[2].set_title("Gradient Boosting")
axs[2].set_xlabel("Predicted Labels")
axs[2].set_ylabel("True Labels")

# Adjust the layout and spacing
plt.tight_layout()
plt.show()
```



A plot of ROC and AUC

```

In [98]: # Initialize dictionaries to store the evaluation metrics for each model
accuracy_scores = {}
precision_scores = {}
recall_scores = {}
f1_scores = {}

# Assuming "models" is a dictionary of your tuned models (Gradient Boosting, Random Forest, etc.)
for model_name, model in models.items():
    # Make predictions on the test set
    y_pred = model.predict(X_test)

    # Evaluate the model
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    # Store the evaluation metrics in the dictionaries
    accuracy_scores[model_name] = accuracy
    precision_scores[model_name] = precision
    recall_scores[model_name] = recall
    f1_scores[model_name] = f1

# Create subplots for the bar plots
fig, axs = plt.subplots(2, 2, figsize=(18, 12))

# Visualize the evaluation metrics using bar plots
axs[0, 0].bar(accuracy_scores.keys(), accuracy_scores.values())
axs[0, 0].set_ylim(0, 1.0)
axs[0, 0].set_title("Accuracy Scores")
axs[0, 0].set_ylabel("Accuracy")

axs[0, 1].bar(precision_scores.keys(), precision_scores.values())
axs[0, 1].set_ylim(0, 1.0)
axs[0, 1].set_title("Precision Scores")
axs[0, 1].set_ylabel("Precision")

axs[1, 0].bar(recall_scores.keys(), recall_scores.values())
axs[1, 0].set_ylim(0, 1.0)
axs[1, 0].set_title("Recall Scores")
axs[1, 0].set_xlabel("Model")
axs[1, 0].set_ylabel("Recall")

axs[1, 1].bar(f1_scores.keys(), f1_scores.values())
axs[1, 1].set_ylim(0, 1.0)
axs[1, 1].set_title("F1-scores")
axs[1, 1].set_xlabel("Model")
axs[1, 1].set_ylabel("F1-score")

plt.tight_layout()
plt.show()

# Create a new figure for the ROC curves
plt.figure(figsize=(10, 8))

# Create ROC curves for each model
for model_name, model in models.items():
    y_prob = model.predict_proba(X_test)[:, 1]
    # Create the ROC curve for each model
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, lw=2, label=f'{model_name} (AUC: {roc_auc})')
    plt.plot([0, 1], [0, 1], lw=1, label='Random Guess')

```


```
    tpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f"{model_name} (AUC = {roc_auc:.2f})")


# Plot the diagonal line, which represents a random classifier
plt.plot([0, 1], [0, 1], color='grey', lw=2, linestyle='--')

# Set labels and title
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')

# Show the Legend
plt.legend(loc='lower right')

# Display the plot
plt.show()
```





ROC curves takeaways

The ROC (Receiver Operating Characteristic) curve evaluates a model's ability to distinguish between churners and non-churners. The AUC (Area Under the Curve) value quantifies the overall performance, where 1.0 represents a perfect model and 0.5 represents a random guess.

Gradient Boosting (AUC = 0.94) → Best Model

The green curve (Gradient Boosting) stays closest to the top-left corner. AUC of 0.94 suggests that this model is highly effective in distinguishing churners from non-churners.

Business Implication: This model can be confidently used to predict and take action on potential churners.

Random Forest (AUC = 0.93) → Second-Best Model

The orange curve is slightly below the Gradient Boosting curve. AUC of 0.93 indicates that Random Forest is also a strong performer.

Business Insight: If computational efficiency or interpretability is a concern, this model is a great alternative.

Logistic Regression (AUC = 0.85) → Lowest Performance

The blue curve is farther from the top-left corner. AUC of 0.85 is still decent but lower than tree-

based models.

Business Insight: Logistic Regression struggles with non-linear relationships, making it less suitable for this dataset.

```
In [99]: # Results before tuning
results_df
```

```
Out[99]:
```

| | | Accuracy | Precision | Recall | F1-score |
|----------------------------|--------------|----------|-----------|----------|----------|
| Logistic Regression | Test | 0.767722 | 0.325301 | 0.562500 | 0.412214 |
| | Train | 0.770857 | 0.327613 | 0.544041 | 0.408958 |
| Random Forest | Test | 0.904977 | 0.714286 | 0.572917 | 0.635838 |
| | Train | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| Gradient Boosting | Test | 0.960784 | 0.897727 | 0.822917 | 0.858696 |
| | Train | 0.975840 | 0.976331 | 0.854922 | 0.911602 |

```
In [100... evaluate_model(best_gradient_boosting_model, X_train, y_train, X_test, y_test)
```

Training Data - Accuracy: 1.0000, Precision: 1.0000, Recall: 1.0000, F1-score: 1.0000
Test Data - Accuracy: 0.9713, Precision: 0.9873, Recall: 0.8125, F1-score: 0.8914

Summary of the performance of the various models

1. Baseline Model: Decision Tree

Before Tuning:

Overfitting is evident since the training accuracy is 100%, but test accuracy drops to 94.72%. The recall of 81.25% on the test set means the model is capturing most churned cases but still missing some. Precision of 82.11% shows that when the model predicts churn, it is correct most of the time.

After Tuning:

Training accuracy slightly drops to 98.48%, reducing overfitting. Test accuracy slightly decreases to 93.67%, but recall improves to 86.46%, meaning it captures more churned customers. However, precision drops to 74.11%, meaning the model makes more false positives. Trade-off: The model favors identifying more churn cases but at the cost of slightly more false positives.

2. Logistic Regression

Before Tuning:

Test accuracy is 83.41%, significantly lower than tree-based models. Recall is 64.58%, meaning it captures churn cases moderately well, but precision is 44.93%, so many non-churned customers are misclassified as churned.

After Tuning:

Test accuracy improves to 87.33%, but recall drops to 31.25%, meaning it now struggles to identify churned customers. Precision increases to 62.50%, meaning when it predicts churn, it is more likely correct. Trade-off: The tuned model is more conservative in predicting churn, leading to fewer false positives but missing more actual churn cases.

3. Random Forest

Before Tuning:

Training accuracy is 100% (overfitting). Test accuracy is 93.36%, recall is 80.21%, and precision is 75.49% (balanced performance).

After Tuning:

Test accuracy remains 93.36%, but recall drops to 54.17%, meaning the model misses more churned customers. Precision jumps to 100%, meaning when it predicts churn, it is always correct. Trade-off: The model is highly conservative—preferring not to predict churn unless it is absolutely sure, which increases false negatives.

4. Gradient Boosting

Before Tuning:

Test accuracy is 97.59%, better than other models. Precision is 96.51%, recall is 86.46%—this means it predicts churn cases accurately while still capturing most of them. The best balance between precision and recall compared to the other models.

After Tuning:

Training accuracy becomes 100% (potential overfitting). Test accuracy is 96.68%, slightly dropping, but still better than other models. Precision is 95.12%, recall is 81.25%—meaning the model is still strong in capturing churned customers and avoiding false positives. Trade-off: Gradient boosting provides the best balance between precision and recall while maintaining high accuracy.

Overall, Gradient Boosting demonstrates the best performance among the three models, achieving high accuracy and balanced precision-recall trade-off on the test set.

Overall Insights and Best Model Choice

Gradient Boosting (After Tuning) Performs Best

Highest test accuracy (96.68%) with good recall (81.25%) and excellent precision (95.12%). Outperforms other models in balancing recall and precision. Slight risk of overfitting, but generalizes well to test data.

Decision Tree (After Tuning) is a Good Alternative

Good recall (86.46%) but lower precision (74.11%) compared to gradient boosting. If high recall is the priority (capturing all churn cases), this model is preferable.

Random Forest (After Tuning) is Too Conservative

100% precision but only 54.17% recall—it captures very few churn cases. Would only be useful if the cost of false positives is extremely high.

Logistic Regression is the Weakest Model

Even after tuning, recall drops significantly (31.25%), making it unreliable for churn prediction.

Final Recommendation

For the best balance between accuracy, recall, and precision, Tuned Gradient Boosting is the best model for predicting churn in this dataset.

After tuning, the model is slightly more conservative, meaning it avoids over-predicting churn (false positives). This makes it more reliable in real-world deployment, where predicting too many non-churners as churners can have unnecessary business costs (e.g., offering retention incentives to customers who were never going to leave).

Feature Importance


In [101...

```
n_features = best_clf.n_features_in_
threshold = 0.01 # Set the threshold for including features with importance

# Filter the features based on the importance threshold
important_features = [feature for feature, importance in zip(X.columns, best_clf.feature_importances_)]

# Create a subset of feature importances for the important features
importances_subset = [importance for importance in best_clf.feature_importances_ if importance > threshold]

plt.figure(figsize=(30, 20))
plt.barh(range(len(important_features)), importances_subset)
plt.yticks(range(len(important_features)), important_features, fontsize=20)
plt.xlabel('Feature Importance', fontsize=20)
plt.ylabel('Feature', fontsize=30)
plt.title(f'FSM Feature Importance (Threshold: >{threshold})', fontsize=20)
plt.tight_layout()
plt.show()
```



Feature Importance Analysis

From the above, we can break down the key findings:

Key Features and Their Importance

Customer Service Calls (Most Important Feature)

This is the most critical factor influencing churn. A higher number of customer service calls is strongly associated with churn, suggesting that dissatisfied customers frequently contact support before leaving.

Business Insight: The company should analyze customer service interactions, improve support quality, and resolve complaints efficiently to reduce churn.

Total Charge

This feature likely includes total charges across different time periods (day, evening, night, international). Customers with higher charges may feel dissatisfied with their billing or find cheaper alternatives.

Business Insight: Offering better pricing plans, discounts, or loyalty programs can help retain high-paying customers.

International Plan (Yes/No)

Customers subscribed to an international plan appear to have a higher likelihood of churning. This could mean that the plan is not meeting their expectations, or they are finding better deals elsewhere.

Business Insight: The company should assess the competitiveness of its international plans and offer better rates or bundled services.

Voice Mail Message to Call Ratio

A higher ratio may indicate reliance on voicemail rather than direct communication. Possible interpretation: Customers who rely heavily on voicemail may have different communication needs that are not being met.

Business Insight: Investigate how voicemail users interact with the service and whether alternative features like unlimited calling or messaging can improve retention.

Total Talk Time

More time spent on calls could indicate a heavy user segment. These users may be prone to switching to competitors offering better deals.

Business Insight: Consider targeting high-usage customers with retention offers or exclusive benefits.

Total International Charge

Higher international charges may be a reason for churn, as users seek cheaper alternatives.

Business Insight: Offer international calling discounts or partnerships with global carriers.

Total International Calls & Total Day Calls

Both features have lower importance but still contribute to the prediction. International calls might indicate a specific customer segment that can be targeted for tailored offers.

Business Insight: Market specialized international packages to frequent international callers.

Account Length

Surprisingly, this has relatively low importance. It suggests that tenure alone is not a strong predictor of churn—both new and long-term customers can churn.

Business Insight: Focus more on behavioral patterns than account age when designing retention strategies.

Recommendations

Based on the findings, I recommend the following:

1. Improve Customer Support:

Since customer service calls are the biggest churn factor, identify common complaints and resolve them proactively. Use AI-driven support or faster resolution times to improve customer satisfaction.

2. Target High-Charge Customers:

Customers with high total charges should receive loyalty rewards, exclusive discounts, or personalized offers to reduce price sensitivity.

3. Re-Evaluate International Plans:

Since international plan users are more likely to churn, conduct a survey to understand why. Offer flexible international calling packages or discounts.

4. Enhance Retention Efforts Based on Usage:

Customers with high talk time and high call volumes should be offered personalized plans to match their needs.

5. Monitor Voicemail Users:

Users with high voicemail usage may prefer messaging or digital alternatives. Consider promoting chat-based or unlimited text/calling plans.